



AutoMLOps

From Notebooks to Pipelines in Minutes

Implementation Guide v1.0.5

srastatter@, tonydiloreto@, allegranoto@

February 2023

The diagram illustrates the architecture of AutoMLOps > CI/CD & Orchestration on Google Cloud Platform. It shows the flow from code generation to model training and deployment.

Components and Flow:

- Source Code (CSR):** Cloud Storage Repository.
- Run Build on Push:** Cloud Build Trigger.
- Build Base Image:** Cloud Build.
- Build Runner SVC:** Cloud Build.
- Deploy Runner SVC:** Cloud Build.
- Job Queueing SVC:** Cloud Tasks.
- Packages:** Artifact Registry.
- Run PipelineJob:** Cloud Run.
- Scheduled Run:** Cloud Scheduler.
- Training Job:** Vertex AI Pipelines.
- Generate/Run Code:** AutoMLOps.go.
- Generate Code:** AutoMLOps.generate.
- Call AutoMLOps:** Vertex AI Workbench.

Flow Description:

- Call AutoMLOps (Vertex AI Workbench)** triggers **Generate Code (AutoMLOps.generate)**.
- Generate Code (AutoMLOps.generate)** triggers **Generate/Run Code (AutoMLOps.go)**.
- Generate/Run Code (AutoMLOps.go)** triggers **Source Code (CSR)**.
- Source Code (CSR)** triggers **Run Build on Push (Cloud Build Trigger)**.
- Run Build on Push (Cloud Build Trigger)** triggers **Build Base Image (Cloud Build)**, **Build Runner SVC (Cloud Build)**, **Deploy Runner SVC (Cloud Build)**, and **Job Queueing SVC (Cloud Tasks)**.
- Build Base Image (Cloud Build)** and **Build Runner SVC (Cloud Build)** trigger **Packages (Artifact Registry)**.
- Packages (Artifact Registry)** triggers **Run PipelineJob (Cloud Run)**.
- Deploy Runner SVC (Cloud Build)** triggers **Run PipelineJob (Cloud Run)**.
- Job Queueing SVC (Cloud Tasks)** **<Invokes>** **Run PipelineJob (Cloud Run)**.
- Scheduled Run (Cloud Scheduler)** triggers **Run PipelineJob (Cloud Run)**.
- Run PipelineJob (Cloud Run)** triggers **Training Job (Vertex AI Pipelines)**.

Architecture: AutoMLOps > CI/CD & Orchestration

Set Up



Prerequisites / Assumptions

The prerequisites for use of AutoMLOps are as follows:

- Jupyter (or Jupyter-compatible) notebook environment
- [Notebooks API](#) is enabled
- Python version ≥ 3.0 and ≤ 3.10
- [Google Cloud SDK 407.0.0](#)
- [gcloud beta 2022.10.21](#)
- git is installed and logged-in

```
git config --global user.email "you@example.com"  
git config --global user.name "Your Name"
```

- [Application Default Credentials \(ADC\)](#) are set up, which can be done through the following commands:

```
gcloud auth application-default login  
gcloud config set account <account@example.com>
```

Set Up AutoMLOps Package

1. Install the AutoMLOps package:

```
pip install google-cloud-automlops
```

2. Open a notebook and import the AutoMLOps package:

```
from AutoMLOps import AutoMLOps
```

3. Decide whether to use Kubeflow definitions or Python definitions

Using Python Components (no Kubeflow)



AutoMLOps Python Process

1. Define import code cells

```
%%define_imports
import json
import pandas as pd
from google.cloud import aiplatform
from google.cloud import aiplatform_v1
from google.cloud import bigquery
...
```

AutoMLOps Python Process

2. Define component code cells

```
%%define_component
AutoMLOps.makeComponent(
    name = "create_dataset",
    description = "Loads data from BQ and writes a dataframe as a csv to GCS.", # optional
    params = [
        {"name": "bq_table", "type": str, # descriptions are optional
        {"name": "data_path", "type": str, "description": "GS location where the training data is written."},
        {"name": "project_id", "type": str, "description": "Project_id."}
    ]
)
```


AutoMLOps Python Process

3. Define pipeline code cell

```
AutoMLOps.makePipeline(  
    name = "training-pipeline",  
    description = "description", # optional  
    params = [  
        {"name": "bq_table", "type": str}, # descriptions are optional  
        {"name": "model_directory", "type": str, "description": "Description."},  
        {"name": "data_path", "type": str, "description": "Description."},  
        {"name": "project_id", "type": str, "description": "Description."},  
        {"name": "region", "type": str, "description": "Description."}  
    ],  
    pipeline = [{  
        "component_name": "create_dataset", "param_mapping": [  
            ("bq_table", "bq_table"), # (component_param, pipeline_param)  
            ("data_path", "data_path"),  
            ("project_id", "project_id")  
        ]  
    },  
    {  
        "component_name": "train_model", "param_mapping": [  
            ("model_directory", "model_directory"),  
            ("data_path", "data_path")  
        ]  
    },  
    {  
        "component_name": "deploy_model", "param_mapping": [  
            ("model_directory", "model_directory"),  
            ("project_id", "project_id"),  
            ("region", "region")  
        ]  
    }  
    ]  
)
```

AutoMLOps Python Process

4. Define the pipeline parameters dictionary

```
pipeline_params = {  
    "bq_table": f"{PROJECT_ID}.test_dataset.dry-beans",  
    "model_directory": f"gs://{PROJECT_ID}-bucket/trained_models/{datetime.datetime.now()}",  
    "data_path": f"gs://{PROJECT_ID}-bucket/data",  
    "project_id": f"{PROJECT_ID}",  
    "region": "us-central1"  
}
```

AutoMLOps Python Process

5. Call `AutoMLOps.generate()` to create the resources and repository
Or `AutoMLOps.go()` to call generate in addition to building/submitting the pipeline job

```
AutoMLOps.generate(project_id = PROJECT_ID,  
                    pipeline_params = pipeline_params,  
                    use_kfp_spec = False,  
                    run_local = True)  
  
AutoMLOps.go(project_id = PROJECT_ID,  
              pipeline_params = pipeline_params,  
              use_kfp_spec = False,  
              run_local = True)
```

Set `use_kfp_spec=False` when using an AutoMLOps Python defined pipeline

Set `run_local=False` if you want to generate and use CI/CD features

Using Kubeflow Components



AutoMLOps Kubeflow Process

1. Define your components using KFP

```
@component(  
    packages_to_install = [  
        "google-cloud-bigquery",  
        "pandas",  
        "pyarrow",  
        "db_dtypes"  
    ],  
    base_image = "python:3.9",  
    output_component_file = f"{AutoMLOps.OUTPUT_DIR}/create_dataset.yaml"  
)  
def create_dataset(  
    bq_table: str,  
    output_data_path: OutputPath("Dataset"),  
    project: str  
):  
    from google.cloud import bigquery  
    ...
```

AutoMLOps Kubeflow Process

2. If using Kubeflow defs, define your pipeline using KFP and AutoMLOps:

```
%%define_kfp_pipeline

@dsl.pipeline(name = 'training-pipeline')
def pipeline(bq_table: str,
             output_model_directory: str,
             project: str,
             region: str,
             ):

    dataset_task = create_dataset(
        bq_table = bq_table,
        project = project)

    model_task = train_model(
        output_model_directory = output_model_directory,
        dataset = dataset_task.output)

    deploy_task = deploy_model(
        model = model_task.outputs["model"],
        project = project,
        region = region)

    ...
```

AutoMLOps Kubeflow Process

3. Define the pipeline parameters dictionary

```
pipeline_params = {  
    "bq_table": f"{PROJECT_ID}.test_dataset.dry-beans",  
    "output_model_directory": f"gs://{PROJECT_ID}-bucket/trained_models/{datetime.datetime.now()}",  
    "project": f"{PROJECT_ID}",  
    "region": "us-central1"  
}
```

AutoMLOps Kubeflow Process

4. Call `AutoMLOps.generate()` to create the resources and repository
Or `AutoMLOps.go()` to call generate in addition to building/submitting the pipeline job

```
AutoMLOps.generate(project_id = PROJECT_ID,  
                   pipeline_params = pipeline_params,  
                   use_kfp_spec = True,  
                   run_local = False)  
  
AutoMLOps.go(project_id = PROJECT_ID,  
              pipeline_params = pipeline_params,  
              use_kfp_spec = True,  
              run_local = False)
```

Set `use_kfp_spec=True` when using a Kubeflow defined pipeline

Set `run_local=False` if you want to generate and use CI/CD features

Customizations



AutoMLOps Defaults

There are a number of custom parameters that can be configured. To the left is a list of the default parameters:

```
AutoMLOps.go(project_id=PROJECT_ID, # required
               pipeline_params=pipeline_params, # required
               af_registry_location='us-central1', # default
               af_registry_name='vertex-mlops-af', # default
               cb_trigger_location='us-central1', # default
               cb_trigger_name='automlops-trigger', # default
               cloud_run_location='us-central1', # default
               cloud_run_name='run-pipeline', # default
               cloud_tasks_queue_location='us-central1', # default
               cloud_tasks_queue_name='queueing-svc', # default
               csr_branch_name='automlops', # default
               csr_name='AutoMLOps-repo', # default
               custom_training_job_specs=None, # default
               gs_bucket_location='us-central1', # default
               gs_bucket_name=None, # default
               pipeline_runner_sa=None, # default
               run_local=True, # default
               schedule_location='us-central1', # default
               schedule_name='AutoMLOps-schedule', # default
               schedule_pattern='No Schedule Specified', # default
               use_kfp_spec=False, # default
               vpc_connector='No VPC Specified' # default)
```

AutoMLOps Defaults

- ``project_id``: The project ID.
- ``pipeline_params``: Dictionary containing runtime pipeline parameters.
- ``af_registry_location``: Region of the Artifact Registry.
- ``af_registry_name``: Artifact Registry name where components are stored.
- ``cb_trigger_location``: The location of the cloudbuild trigger.
- ``cb_trigger_name``: The name of the cloudbuild trigger.
- ``cloud_run_location``: The location of the cloud runner service.
- ``cloud_run_name``: The name of the cloud runner service.
- ``cloud_tasks_queue_location``: The location of the cloud tasks queue.
- ``cloud_tasks_queue_name``: The name of the cloud tasks queue.
- ``csr_branch_name``: The name of the csr branch to push to to trigger cb job.
- ``csr_name``: The name of the cloud source repo to use.
- ``custom_training_job_specs``: Specifies the specs to run the training job with.
- ``gs_bucket_location``: Region of the GS bucket.
- ``gs_bucket_name``: GS bucket name where pipeline run metadata is stored.
- ``pipeline_runner_sa``: Service Account to runner PipelineJobs.
- ``run_local``: Flag that determines whether to use Cloud Run CI/CD.
- ``schedule_location``: The location of the scheduler resource.
- ``schedule_name``: The name of the scheduler resource.
- ``schedule_pattern``: Cron formatted value used to create a Scheduled retrain job.
- ``use_kfp_spec``: Flag that specifies whether you are using Kubeflow spec or not.
- ``vpc_connector``: The name of the vpc connector to use.

AutoMLOps Set Scheduled Run

Use the *schedule_pattern* parameter to specify a cron job schedule to run the pipeline job on a recurring basis.

The *run_local* must be set to *False* to make use of this feature.

```
AutoMLOps.generate(project_id = PROJECT_ID,  
                    pipeline_params = pipeline_params,  
                    use_kfp_spec = False,  
                    run_local = False,  
                    schedule_pattern = '0 */12 * * *')
```

The above example will rerun the pipeline every 12 hours.

AutoMLOps Set Pipeline Compute Resources

Use the `custom_training_job_specs` parameter to specify resources for any custom component in the pipeline. The `component_spec` must match exactly the name of the custom component.

```
AutoMLOps.generate(project_id = PROJECT_ID,  
                    pipeline_params = pipeline_params,  
                    use_kfp_spec = False,  
                    run_local = False,  
                    custom_training_job_specs = [{  
                        'component_spec': 'train_model',  
                        'display_name': 'train-model-accelerated',  
                        'machine_type': 'a2-highgpu-1g',  
                        'accelerator_type': 'NVIDIA_TESLA_A100',  
                        'accelerator_count': '1'  
                    }])
```

The example above uses a GPU for accelerated training. See [Machine types](#) and [GPUs](#) for more info.

The `custom_training_job_specs` parameter takes in any key-value pair available under [google_cloud_pipeline_components.v1.custom_job.create_custom_training_job_op_from_component](#)

AutoMLOps VPC Connector

Use the `vpc_connector` parameter to specify a vpc connector.

```
AutoMLOps.generate(project_id = PROJECT_ID,  
                    pipeline_params = pipeline_params,  
                    use_kfp_spec = False,  
                    run_local = False,  
                    vpc_connector = 'example-vpc')
```

Behind the Scenes



Cloud Resources

When the pipeline is run, the following resources are created to complete the MLOps pipeline:

1. AutoMLOps codebase
2. Artifact Registry
3. GS Bucket
4. Pipeline Runner Service Account
5. Cloud Source Repository (turns the notebooks working directory into a CSR)
6. Cloud Build Trigger
7. Cloud Runner Service
8. Cloud Scheduler
9. Cloud Tasks queue

APIs

When the pipeline is run, the following APIs are enabled:

1. `cloudresourcemanager.googleapis.com`
2. `aiplatform.googleapis.com`
3. `artifactregistry.googleapis.com`
4. `cloudbuild.googleapis.com`
5. `cloudscheduler.googleapis.com`
6. `cloudtasks.googleapis.com`
7. `compute.googleapis.com`
8. `iam.googleapis.com`
9. `iamcredentials.googleapis.com`
10. `ml.googleapis.com`
11. `run.googleapis.com`
12. `storage.googleapis.com`
13. `sourcerepo.googleapis.com`

IAM Access

When the pipeline is run, the following IAM access roles are updated:

1. **Pipeline Runner Service Account** (created if it does exist, defaults to: *vertex-pipelines@<PROJECT_ID>.iam.gserviceaccount.com*).

Roles added:

- roles/aiplatform.user
- roles/artifactregistry.reader
- roles/bigquery.user
- roles/bigquery.dataEditor
- roles/iam.serviceAccountUser
- roles/storage.admin
- roles/run.admin

2. **Cloudbuild Default Service Account** (<PROJECT_NUMBER>@cloudbuild.gserviceaccount.com).

Roles added:

- roles/run.admin
- roles/iam.serviceAccountUser
- roles/cloudtasks.enqueueur
- roles/cloudscheduler.admin

Package Dependencies

When using AutoMLOps, the following package versions are used:

1. autoflake==2.0.0,
2. docopt==0.6.2,
3. ipython==7.34.0,
4. pipreqs==0.4.11,
5. pyflakes==3.0.1,
6. PyYAML==5.4.1,
7. yarg==0.1.9